

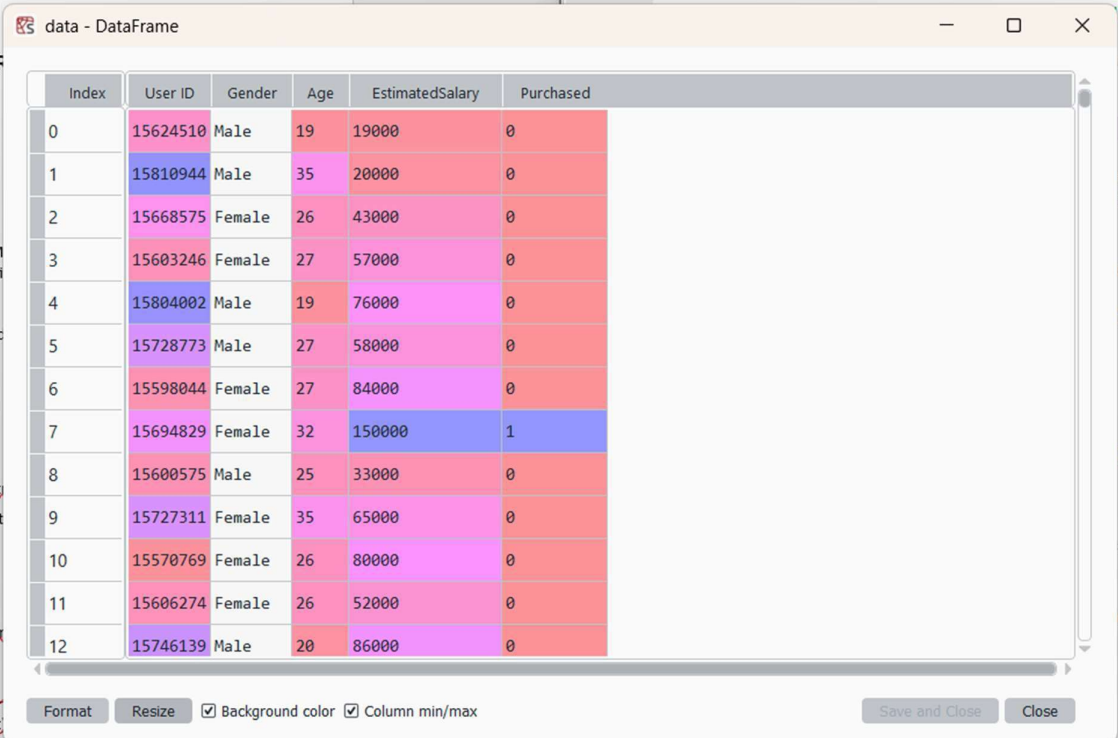
KNN-Classifier

#lets import libraries

```
import pandas as pd
```

#lets read the dataset

```
data=pd.read_csv(r"C:\Users\TharunMahendra\NIT\6.Algorithms\2.Classification\ClassificationData.csv")
```



Index	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1
8	15600575	Male	25	33000	0
9	15727311	Female	35	65000	0
10	15570769	Female	26	80000	0
11	15606274	Female	26	52000	0
12	15746139	Male	20	86000	0

#lets divide them into dependent & independent

```
x=data.iloc[:,2:4] #age,salary
```

```
y=data.iloc[:, -1] #purchased
```

x - DataFrame			y - Series	
Index	Age	EstimatedSalary	Index	Purchased
0	19	19000	0	0
1	35	20000	1	0
2	26	43000	2	0
3	27	57000	3	0
4	19	76000	4	0
5	27	58000	5	0
6	27	84000	6	0
7	32	150000	7	1
8	25	33000	8	0
9	35	65000	9	0
10	26	80000	10	0
11	26	52000	11	0
12	20	86000	12	0

#splitting data

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,
random_state=0)
```

x_test - DataFrame			x_train - DataFrame			y_test - Series		y_train - Series	
Index	Age	EstimatedSalary	Index	Age	EstimatedSalary	Index	Purchased	Index	Purchased
132	30	87000	336	58	144000	132	0	336	1
309	38	50000	64	59	83000	309	0	64	0
341	35	75000	55	24	55000	341	0	55	0
196	30	79000	106	26	35000	196	0	106	0
246	35	50000	300	58	38000	246	0	300	1
60	27	20000	229	42	80000	60	0	229	1
155	31	15000	122	40	75000	155	0	122	0
261	36	144000	373	59	130000	261	1	373	1
141	18	68000	395	46	41000	141	0	395	1
214	47	43000	325	41	60000	214	0	325	0
37	30	49000	380	42	64000	37	0	380	0
134	28	55000	253	37	146000	134	0	253	1
113	37	55000	56	23	48000	113	0	56	0

#feature scaling

from sklearn.preprocessing import **StandardScaler** #range between-> -3to3

featurescaling=StandardScaler()

x_train=featurescaling.fit_transform(x_train)

x_test=featurescaling.transform(x_test)

x_test - NumPy object array			x_train - NumPy object array		
	0	1		0	1
0	-0.798951	0.494608	0	1.92295	2.14602
1	-0.0212649	-0.577359	1	2.02016	0.378719
2	-0.312897	0.146943	2	-1.38222	-0.432499
3	-0.798951	0.262831	3	-1.18779	-1.01194
4	-0.312897	-0.577359	4	1.92295	-0.925024
5	-1.09058	-1.44652	5	0.367578	0.291803
6	-0.70174	-1.59138	6	0.173157	0.146943
7	-0.215686	2.14602	7	2.02016	1.74041
8	-1.96548	-0.0558618	8	0.756421	-0.838108
9	0.853632	-0.780164	9	0.270367	-0.287638
10	-0.798951	-0.606331	10	0.367578	-0.17175
11	-0.993372	-0.432499	11	-0.118476	2.20396
12	-0.118476	-0.432499	12	-1.47943	-0.635303

#model building

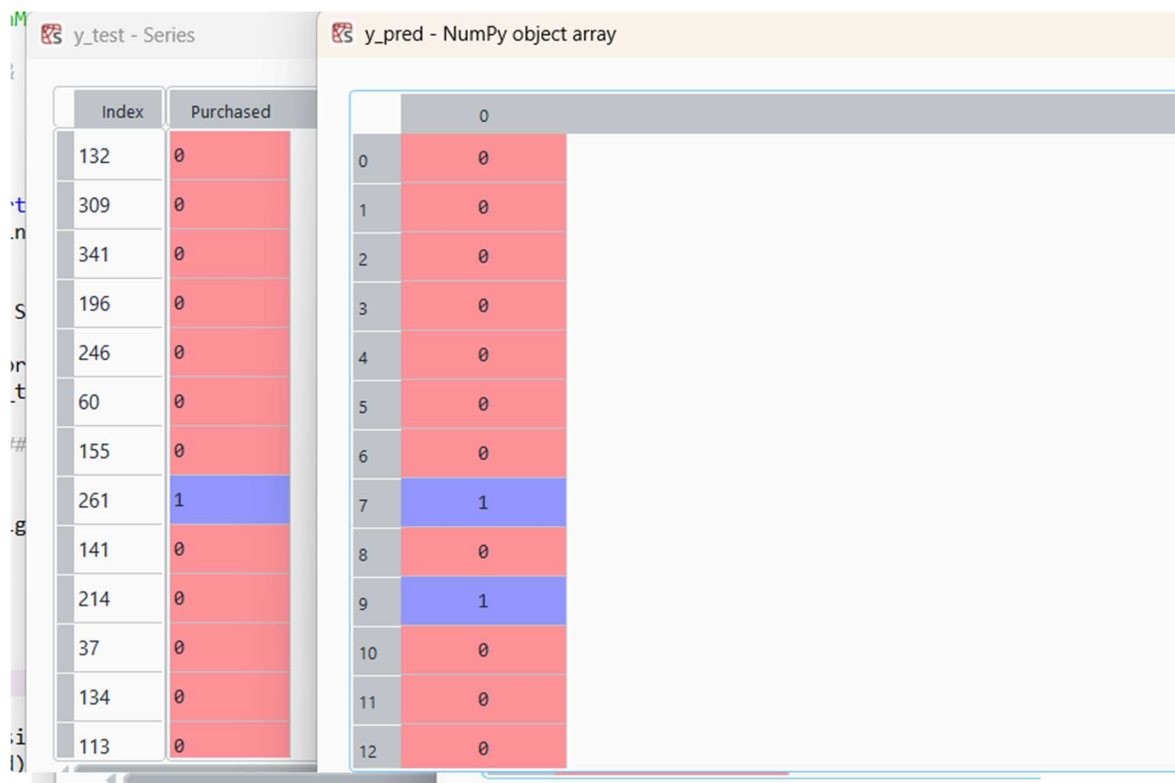
```
from sklearn.neighbors import KNeighborsClassifier
```

```
model=KNeighborsClassifier()
```

```
model.fit(x_train,y_train)
```

#prediction

```
y_pred=model.predict(x_test)
```



The screenshot displays two data structures in a Jupyter Notebook interface. On the left, a pandas Series named 'y_test' is shown with a 'Purchased' column. It contains 15 rows of data, where most values are 0 (red) and one value is 1 (blue) at index 261. On the right, a NumPy array named 'y_pred' is shown, which is a 1D array of integers. It contains 13 rows of data, where most values are 0 (red) and two values are 1 (blue) at indices 7 and 9.

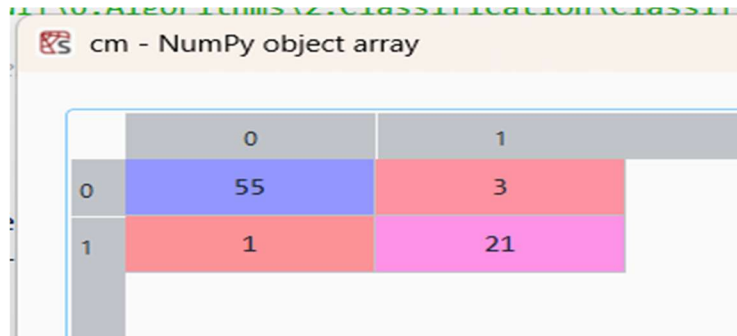
Index	Purchased
132	0
309	0
341	0
196	0
246	0
60	0
155	0
261	1
141	0
214	0
37	0
134	0
113	0

	0
0	0
1	0
2	0
3	0
4	0
5	0
6	0
7	1
8	0
9	1
10	0
11	0
12	0

#confusion Matrix

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test, y_pred)
```



cm - NumPy object array

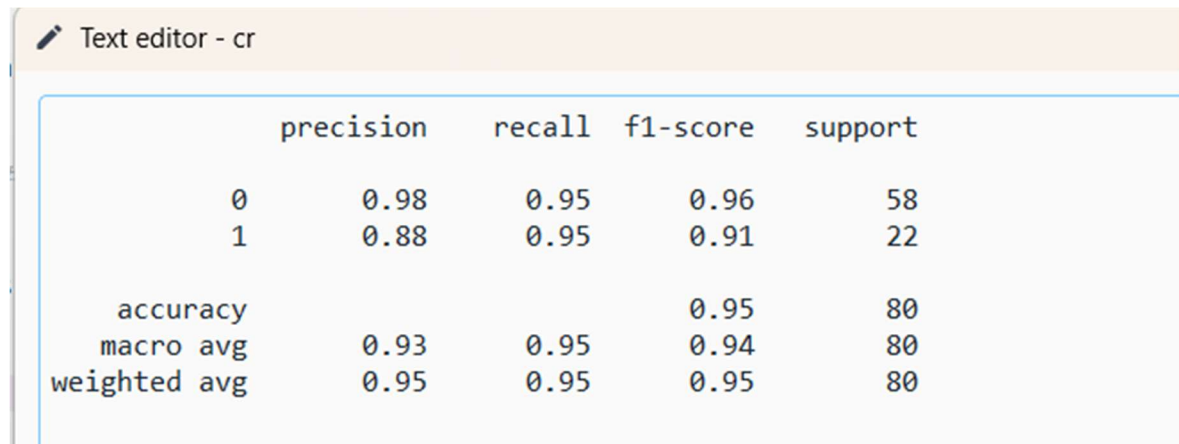
	0	1
0	55	3
1	1	21

```
from sklearn.metrics import accuracy_score
```

```
ac=accuracy_score(y_test,y_pred) ->0.95
```

```
from sklearn.metrics import classification_report
```

```
cr=classification_report(y_test,y_pred)
```



Text editor - cr

	precision	recall	f1-score	support
0	0.98	0.95	0.96	58
1	0.88	0.95	0.91	22
accuracy			0.95	80
macro avg	0.93	0.95	0.94	80
weighted avg	0.95	0.95	0.95	80

```
bias=model.score(x_train,y_train) -> 0.91875
```

```
variance=model.score(x_test,y_test)-> 0.95
```

Visualising the Training set results

```
from matplotlib.colors import ListedColormap

X_set, y_set = x_train, y_train

X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:,
0].max() + 1, step = 0.01),

                     np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1,
step = 0.01))

plt.contourf(X1, X2, model.predict(np.array([X1.ravel(),
X2.ravel()]).T).reshape(X1.shape),

             alpha = 0.75, cmap = ListedColormap(('red', 'green')))

plt.xlim(X1.min(), X1.max())

plt.ylim(X2.min(), X2.max())

for i, j in enumerate(np.unique(y_set)):

    plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],

                c = ListedColormap(('red', 'green'))(i), label = j)


plt.title('K-NN (Training set)')

plt.xlabel('Age')

plt.ylabel('Estimated Salary')

plt.legend()

plt.show()
```



Visualising the Test set results

```
from matplotlib.colors import ListedColormap
```

```
X_set, y_set = x_test, y_test
```

```
X1, X2 = np.meshgrid(np.arange(start = X_set[:, 0].min() - 1, stop = X_set[:, 0].max() + 1, step = 0.01),
```

```
np.arange(start = X_set[:, 1].min() - 1, stop = X_set[:, 1].max() + 1, step = 0.01))
```

```
plt.contourf(X1, X2, model.predict(np.array([X1.ravel(), X2.ravel()]).T).reshape(X1.shape),
```

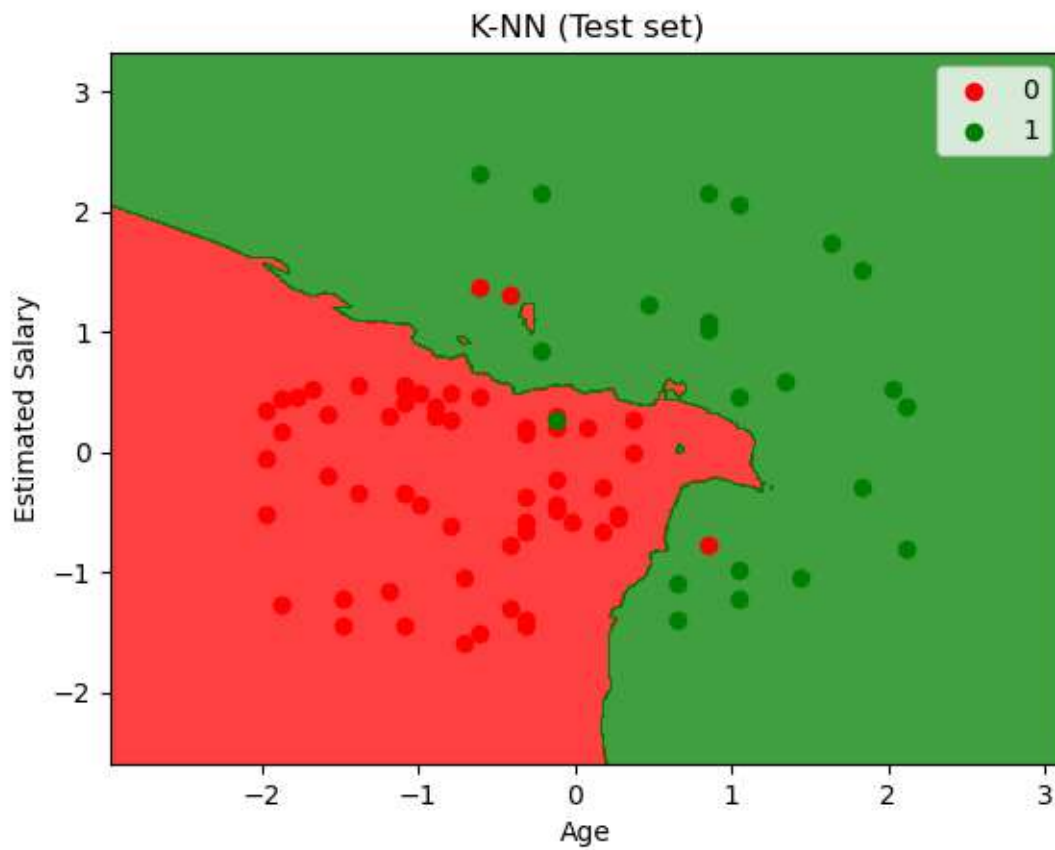
```
alpha = 0.75, cmap = ListedColormap(('red', 'green')))
```

```
plt.xlim(X1.min(), X1.max())
```

```
plt.ylim(X2.min(), X2.max())
```

```
for i, j in enumerate(np.unique(y_set)):
```

```
plt.scatter(X_set[y_set == j, 0], X_set[y_set == j, 1],  
            c = ListedColormap(('red', 'green'))(i), label = j)  
  
plt.title('K-NN (Test set)')  
  
plt.xlabel('Age')  
  
plt.ylabel('Estimated Salary')  
  
plt.legend()  
  
plt.show()
```



Deployment Code

importing libraries

```
import streamlit as st
```

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics
```

```
import(confusion_matrix,accuracy_score,classification_report,roc_curve,roc_auc_score)
```

```
import seaborn as sns
```

```
st.title("KNN Classifier")
```

Uploading File

```
file=st.file_uploader('Upload Your File for Model Building',type=['csv'])
```

```
if file is not None:
```

```
    # Load
```

```
    data=pd.read_csv(file)
```

```
    st.write('- Preview')
```

```
    st.dataframe(data.head())
```

```
    # FeatureSelection
```

```
    x=data.iloc[:,2:4]
```

```
y=data.iloc[:,-1]
```

```
# SplittingData
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_state=0)
```

```
# feature scaling
```

```
featurescaling=StandardScaler()
```

```
x_train=featurescaling.fit_transform(x_train)
```

```
x_test=featurescaling.transform(x_test)
```

```
# model building
```

```
model=KNeighborsClassifier()
```

```
model.fit(x_train,y_train)
```

```
# prediction
```

```
y_pred=model.predict(x_test)
```

```
y_prob=model.predict_proba(x_test)[:,1]
```

```
# Metrics
```

```
st.subheader("Confusion Matrix")
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
fig_cm, ax = plt.subplots()
```

```
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=ax)
```

```
st.pyplot(fig_cm)
```

```
ac = accuracy_score(y_test, y_pred)
```

```
st.write(f"**Accuracy:** {ac:.2f}")
```

```
st.subheader("Classification Report")
```

```
st.text(classification_report(y_test, y_pred))
```

```
st.write(f"**Training Accuracy (Bias):** {model.score(x_train,  
y_train):.2f}")
```

```
st.write(f"**Testing Accuracy (Variance):** {model.score(x_test,  
y_test):.2f}")
```

ROC Curve and AUC

```
fpr, tpr, _ = roc_curve(y_test, y_prob)
```

```
auc_score = roc_auc_score(y_test, y_prob)
```

```
st.write(f"**AUC Score:** {auc_score:.2f}")
```

```
st.subheader("ROC Curve")
```

```
fig_roc, ax = plt.subplots()
```

```
ax.plot(fpr, tpr, color="blue", label=f"ROC curve (AUC = {auc_score:.2f})")
```

```
ax.plot([0, 1], [0, 1], color="gray", linestyle="--")
```

```
ax.set_xlabel("False Positive Rate")
```

```
ax.set_ylabel("True Positive Rate")
```

```
ax.set_title("ROC Curve")
```

```
ax.legend(loc="lower right")
```

```
st.pyplot(fig_roc)
```


Deploymentcode

localhost:8501

Deploy

KNN Classifier

Upload Your File for Model Building

 Drag and drop file here
Limit 200MB per file • CSV

Browse files

ClassificationData.csv

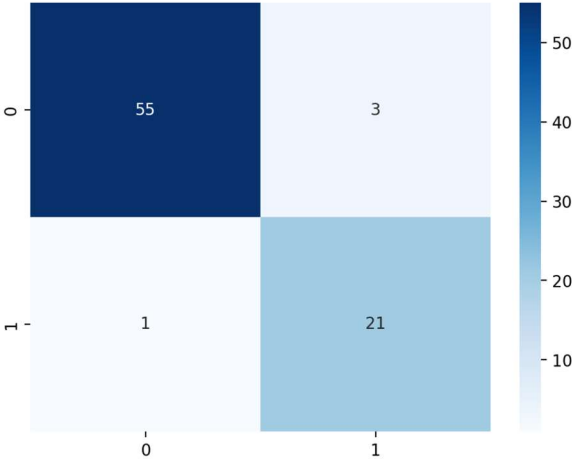
10.7KB

✕

Preview

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15,624,510	Male	19	19,000	0
1	15,810,944	Male	35	20,000	0
2	15,668,575	Female	26	43,000	0
3	15,603,246	Female	27	57,000	0
4	15,804,002	Male	19	76,000	0

Confusion Matrix



	0	1
0	55	3
1	1	21

Accuracy: 0.95

Classification Report

	precision	recall	f1-score	support	
0		0.98	0.95	0.96	58
1		0.88	0.95	0.91	22
accuracy				0.95	80
macro avg		0.93	0.95	0.94	80
weighted avg		0.95	0.95	0.95	80

Training Accuracy (Bias): 0.92

Testing Accuracy (Variance): 0.95

AUC Score: 0.99

