

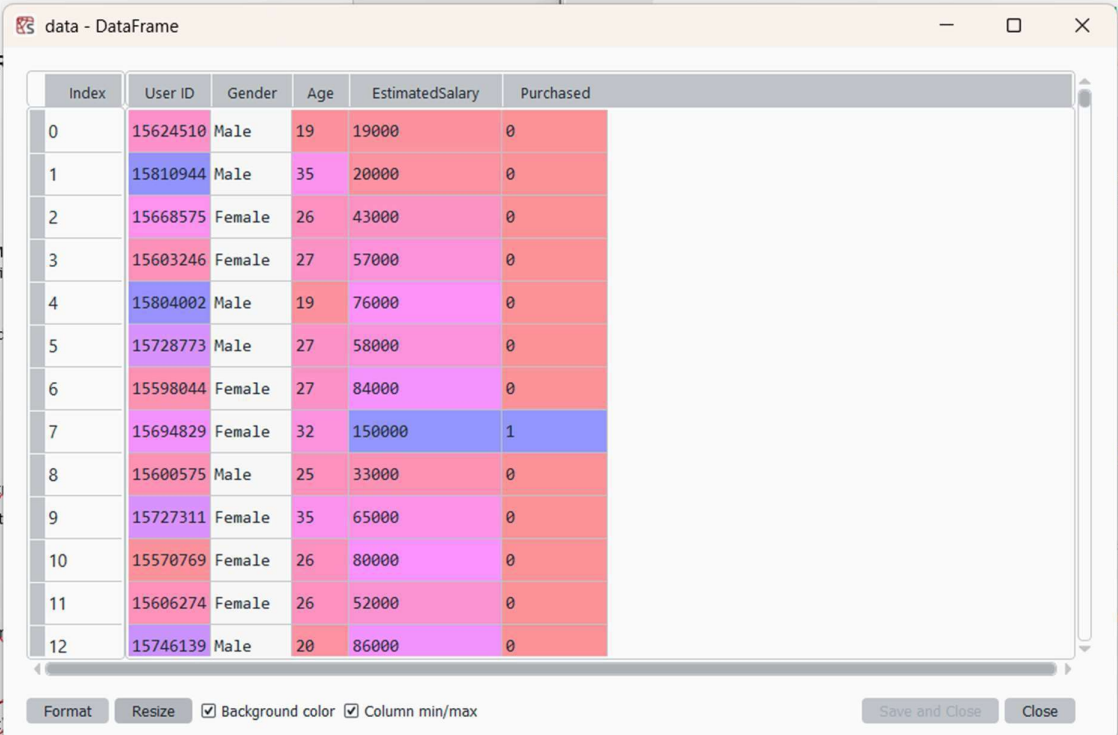
NavieBayes-Classfier

#lets import libraries

```
import pandas as pd
```

#lets read the dataset

```
data=pd.read_csv(r"C:\Users\TharunMahendra\NIT\6.Algorithms\2.Classification\ClassificationData.csv")
```



Index	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1
8	15600575	Male	25	33000	0
9	15727311	Female	35	65000	0
10	15570769	Female	26	80000	0
11	15606274	Female	26	52000	0
12	15746139	Male	20	86000	0

#lets divide them into dependent & independent

```
x=data.iloc[:,2:4] #age,salary
```

```
y=data.iloc[:, -1] #purchased
```

x - DataFrame			y - Series	
Index	Age	EstimatedSalary	Index	Purchased
0	19	19000	0	0
1	35	20000	1	0
2	26	43000	2	0
3	27	57000	3	0
4	19	76000	4	0
5	27	58000	5	0
6	27	84000	6	0
7	32	150000	7	1
8	25	33000	8	0
9	35	65000	9	0
10	26	80000	10	0
11	26	52000	11	0
12	20	86000	12	0

#splitting data

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,
random_state=0)
```

x_test - DataFrame			x_train - DataFrame			y_test - Series		y_train - Series	
Index	Age	EstimatedSalary	Index	Age	EstimatedSalary	Index	Purchased	Index	Purchased
132	30	87000	336	58	144000	132	0	336	1
309	38	50000	64	59	83000	309	0	64	0
341	35	75000	55	24	55000	341	0	55	0
196	30	79000	106	26	35000	196	0	106	0
246	35	50000	300	58	38000	246	0	300	1
60	27	20000	229	42	80000	60	0	229	1
155	31	15000	122	40	75000	155	0	122	0
261	36	144000	373	59	130000	261	1	373	1
141	18	68000	395	46	41000	141	0	395	1
214	47	43000	325	41	60000	214	0	325	0
37	30	49000	380	42	64000	37	0	380	0
134	28	55000	253	37	146000	134	0	253	1
113	37	55000	56	23	48000	113	0	56	0

#feature scaling

from sklearn.preprocessing import **StandardScaler** #range between-> -3to3

featurescaling=StandardScaler()

x_train=featurescaling.fit_transform(x_train)

x_test=featurescaling.transform(x_test)

x_test - NumPy object array			x_train - NumPy object array		
	0	1		0	1
0	-0.798951	0.494608	0	1.92295	2.14602
1	-0.0212649	-0.577359	1	2.02016	0.378719
2	-0.312897	0.146943	2	-1.38222	-0.432499
3	-0.798951	0.262831	3	-1.18779	-1.01194
4	-0.312897	-0.577359	4	1.92295	-0.925024
5	-1.09058	-1.44652	5	0.367578	0.291803
6	-0.70174	-1.59138	6	0.173157	0.146943
7	-0.215686	2.14602	7	2.02016	1.74041
8	-1.96548	-0.0558618	8	0.756421	-0.838108
9	0.853632	-0.780164	9	0.270367	-0.287638
10	-0.798951	-0.606331	10	0.367578	-0.17175
11	-0.993372	-0.432499	11	-0.118476	2.20396
12	-0.118476	-0.432499	12	-1.47943	-0.635303

#model building

BernoulliNB

```
from sklearn.naive_bayes import BernoulliNB, MultinomialNB, GaussianNB
```

```
model=BernoulliNB()
```

```
model.fit(x_train,y_train)
```

#prediction

```
y_pred=model.predict(x_test)
```

y_test - Series		y_pred - NumPy object array	
	Index	Purchased	
	132	0	0
	309	0	1
	341	0	2
	196	0	3
	246	0	4
	60	0	5
	155	0	6
	261	1	7
	141	0	8
	214	0	9
	37	0	10
	134	0	11
	113	0	12

#confusion Matrix

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test, y_pred)
```

cm - NumPy object array

	0	1
0	55	3
1	11	11

```
from sklearn.metrics import accuracy_score
```

```
ac=accuracy_score(y_test,y_pred) -> 0.825
```

```
from sklearn.metrics import classification_report
```

```
cr=classification_report(y_test,y_pred)
```

Text editor - cr

	precision	recall	f1-score	support
0	0.83	0.95	0.89	58
1	0.79	0.50	0.61	22
accuracy			0.82	80
macro avg	0.81	0.72	0.75	80
weighted avg	0.82	0.82	0.81	80

```
bias=model.score(x_train,y_train) -> 0.709375
```

```
variance=model.score(x_test,y_test)-> 0.825
```

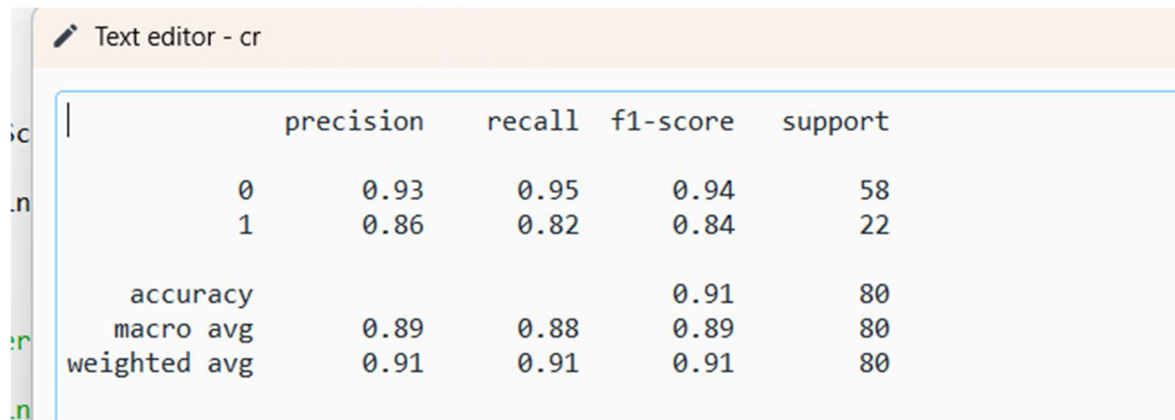
GuassainNB

```
model=GaussianNB()
```

y_test - Series		y_pred - NumPy object array	
Index	Purchased		0
132	0	0	0
309	0	1	0
341	0	2	0
196	0	3	0
246	0	4	0
60	0	5	0
155	0	6	0
261	1	7	1
141	0	8	0
214	0	9	1
37	0	10	0
134	0	11	0
113	0	12	0

cm - NumPy object array		
	0	1
0	55	3
1	4	18

ac=accuracy_score(y_test,y_pred) -> **0.9125**



The screenshot shows a text editor window titled 'Text editor - cr'. Inside, there is a table representing a confusion matrix and summary metrics. The table has five columns: 'precision', 'recall', 'f1-score', and 'support'. The rows are labeled '0', '1', 'accuracy', 'macro avg', and 'weighted avg'. The values are as follows:

	precision	recall	f1-score	support
0	0.93	0.95	0.94	58
1	0.86	0.82	0.84	22
accuracy			0.91	80
macro avg	0.89	0.88	0.89	80
weighted avg	0.91	0.91	0.91	80

bias=model.score(x_train,y_train) -> **0.884375**

variance=model.score(x_test,y_test)-> **0.9125**

MultinomialNB

Fetaure Scaling

```
from sklearn.preprocessing import Normalizer
```

```
featurescaling=Normalizer()
```

```
x_train=featurescaling.fit_transform(x_train)
```

```
x_test=featurescaling.transform(x_test)
```

#we have to use Normalizer for Multionmial Navie-Bayes as Negative Values can't be passed to Multinomial where

Standard Scaler Ranges from -3

```
model=MultinomialNB()
```

```
model.fit(x_train,y_train)
```

y_test - Series		y_pred - NumPy object array	
	Index	Purchased	
	132	0	0
	309	0	0
	341	0	0
	196	0	0
	246	0	0
	60	0	0
	155	0	0
	261	1	0
	141	0	0
	214	0	0
	37	0	0
	134	0	0
	113	0	0

cm - NumPy object array		
	0	1
0	58	0
1	22	0

ac=accuracy_score(y_test,y_pred) -> **0.725**

Text editor - cr

	precision	recall	f1-score	support
0	0.72	1.00	0.84	58
1	0.00	0.00	0.00	22
accuracy			0.72	80
macro avg	0.36	0.50	0.42	80
weighted avg	0.53	0.72	0.61	80

bias=model.score(x_train,y_train) -> **0.621875**

variance=model.score(x_test,y_test)-> **0.72**

Deployment Code

importing libraries

```
import streamlit as st
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, Normalizer
from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB
from sklearn.metrics
import(confusion_matrix, accuracy_score, classification_report, roc_curve, roc_auc_score)
import seaborn as sns
```

```
st.title("NavieBayesClassifier")
```

Uploading File

```
file=st.file_uploader('Upload Your File for Model Building',type=['csv'])
if file is not None:
```

Load

```
data=pd.read_csv(file)
st.write('- Preview')
st.dataframe(data.head())
```

FeatureSelection

```
x=data.iloc[:,2:4]
```

```
y=data.iloc[:,-1]
```

SplittingData

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.20,random_
state=0)
```

feature scaling

```
st.subheader("Feature Scaling")
```

```
scaling_method = st.selectbox("Select ScalingMethod",("StandardScaler",
"Normalizer","None"))
```

```
if scaling_method == "StandardScaler":
```

```
    featurescaling = StandardScaler()
```

```
elif scaling_method == "Normalizer":
```

```
    featurescaling = Normalizer()
```

```
else:
```

```
    featurescaling = None
```

```
if featurescaling is not None:
```

```
    x_train=featurescaling.fit_transform(x_train)
```

```
    x_test=featurescaling.transform(x_test)
```

model building

model selection

```
st.subheader("Select Model")
```

```
model = st.selectbox("Choose a model", ("GaussianNB",
"MultinomialNB", "BernoulliNB"))
```

```
if model == "GaussianNB":
```

```
    model = GaussianNB()
```

```
elif model == "MultinomialNB":
```

```
model = MultinomialNB()
```

```
else:
```

```
model = BernoulliNB()
```

```
model.fit(x_train,y_train)
```

prediction

```
y_pred=model.predict(x_test)
```

```
y_prob=model.predict_proba(x_test)[:,1]
```

Metrics

```
st.subheader("Confusion Matrix")
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
fig_cm, ax = plt.subplots()
```

```
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=ax)
```

```
st.pyplot(fig_cm)
```

```
ac = accuracy_score(y_test, y_pred)
```

```
st.write(f"**Accuracy:** {ac:.2f}")
```

```
st.subheader("Classification Report")
```

```
st.text(classification_report(y_test, y_pred))
```

```
st.write(f"**Training Accuracy (Bias):** {model.score(x_train,  
y_train):.2f}")
```

```
st.write(f"**Testing Accuracy (Variance):** {model.score(x_test,  
y_test):.2f}")
```

ROC Curve and AUC

```
fpr, tpr, _ = roc_curve(y_test, y_prob)
auc_score = roc_auc_score(y_test, y_prob)
st.write(f"**AUC Score:** {auc_score:.2f}")

st.subheader("ROC Curve")
fig_roc, ax = plt.subplots()
ax.plot(fpr, tpr, color="blue", label=f"ROC curve (AUC = {auc_score:.2f})")
ax.plot([0, 1], [0, 1], color="gray", linestyle="--")
ax.set_xlabel("False Positive Rate")
ax.set_ylabel("True Positive Rate")
ax.set_title("ROC Curve")
ax.legend(loc="lower right")
st.pyplot(fig_roc)
```

The screenshot shows a web application titled "NavieBayesClassifier". It has a "Deploy" button in the top right corner. The main content area includes:

- Upload Your File for Model Building:** A section with a "Drag and drop file here" instruction, a "Limit 200MB per file • CSV" note, and a "Browse files" button.
- ClassificationData.csv 10.7KB:** A file upload indicator with a close button (X).
- Preview:** A section showing a table of data.
- Feature Scaling:** A section with a "Select Scaling Method" dropdown menu.

The data preview table is as follows:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15,624,510	Male	19	19,000	0
1	15,810,944	Male	35	20,000	0
2	15,668,575	Female	26	43,000	0
3	15,603,246	Female	27	57,000	0
4	15,804,002	Male	19	76,000	0

The "Feature Scaling" dropdown menu is currently set to "StandardScaler" and shows a list of options: "StandardScaler", "Normalizer", and "None".

Select Model

Choose a model


GaussianNB

GaussianNB

MultinomialNB

BernoulliNB

Confusion Matrix



Accuracy: 0.91

Classification Report


precision	recall	f1-score	support	
0	0.93	0.95	0.94	58
1	0.86	0.82	0.84	22
accuracy			0.91	80
macro avg	0.89	0.88	0.89	80
weighted avg	0.91	0.91	0.91	80

Training Accuracy (Bias): 0.88

Testing Accuracy (Variance): 0.91

AUC Score: 0.98

ROC Curve



Feature Scaling

Select Scaling Method

StandardScaler

Select Model

Choose a model


BernoulliNB

GaussianNB

MultinomialNB

BernoulliNB

Confusion Matrix



Accuracy: 0.82

Classification Report


precision	recall	f1-score	support		
	0	0.83	0.95	0.89	58
	1	0.79	0.50	0.61	22
accuracy				0.82	80
macro avg	0.81	0.72	0.75	80	
weighted avg	0.82	0.82	0.81	80	

Training Accuracy (Bias): 0.71

Testing Accuracy (Variance): 0.82

AUC Score: 0.90

ROC Curve



Feature Scaling

Select Scaling Method

Normalizer

StandardScaler

Normalizer

None

BernoulliNB

Select Model

Choose a model


MultinomialNB

GaussianNB

MultinomialNB

BernoulliNB

Confusion Matrix



Accuracy: 0.72

Classification Report

precision	recall	f1-score	support		
	0	0.72	1.00	0.84	58
	1	0.00	0.00	0.00	22
accuracy				0.72	80
macro avg	0.36	0.50	0.42		80
weighted avg	0.53	0.72	0.61		80

Training Accuracy (Bias): 0.62

Testing Accuracy (Variance): 0.72

AUC Score: 0.57

ROC Curve

